

Motivation

- Computational Efficiency (inkl. Horowitz) !! Why do we do this?
- Improved Generalization (train with low precision (lp) but keep hp weights and use them during test -> better results) Some power gains during training, no power gains during test.
- Energy efficiency during test time if you use lp weights, sacrifice accuracy for power saving
- Examples and related concepts: dropout, ReLu (sparse gradient), influence of feature discretization,

Binarization Schemes

- Deterministic vs. Stochastic

Gradient Estimation

- Path Derivative Gradient Estimators (reparametrization, ST, SlopeAnnealing)
- Score Function based Gradient Estimators (reinforce+many others)
- Expectation Backpropagation

See Gumble Softmax Paper for overview

Algorithms Adam + BN

- General Review
- Illustrate their importance for QNNs
- Implement algorithm to see how it works without Adam and BN?

Full Algorithm

- Try to visualize, illustrate procedure (highlight similarities and differences to conventional NN training and inference)

Experiments

- Describe different NN architectures (DNN, CNN, RNN)
- Describe data sets (MNIST, CIFAR, ImageNet, PennTree)
- binary vs low precision

Results

- Training results
- Performance results (look up older benchmarks to compare loss in in accuracy to)
- find comparisons from other references

Evolution of the QNN paper

Training NN with lp multiplications

- The initial motivation was to reduce the power consumption of multipliers, which are claimed to be the most power hungry.
- No Binarization of anything yet, only lp weights
- A common pattern already shows: lp multipliers + hp accumulators. Parameter updates are also hp.
- Cost of multiplier $O(\text{precision}^2)$ while cost of accumulator $O(\text{precision})$
- Problem with fixed point: low dynamic range \rightarrow dynamic fixed point.
- Evaluation of maxout on MNIST, CIFAR, SVHN with three formats (floating point, fixed point, dynamic fixed point)
- Techniques: Maxout, dropout, momentum, weight decay, dyn. fixed point, update prc vs prop prec.
- Comparison floating point vs. fixed point ?
- Explain demand for sufficient precision of update (compared to Propagations) due to SGD updates
- Half-precision has little to no impact (hp fine-tuning?)
- Couple of plots showing final test error as function of everything
- References to lp network training in the 90ies!
- No evaluation of computational gains, only show how robust NNs are to lp muls.

Binary Connect: Training DNNs with binary weights during propagations

- Emphasis shifts from efficiency to generalization properties (ala dropout)
- Only binarization of weights, activations not yet binary! Backprop still contains muls!
- mention discretization as form of noise which preserves expected value of weights.
- Again: need for hp accumulation (cite several studies incl haml limited precision one), Interesting: reference which says brain synapse precision is around 6-12 bits
- Relations to dropout, dropconnect (!), variational weight noise?
- DropConnect: Only expected value of weights needs hp
- Another reference about hardware cost of add+mul [22]
- Binarization: Deterministic vs. stochastic, Hard sigmoid more efficient!
- Algorithm: Biases not binarized!
- Tricks: Weight clipping, batch norm, adam, normalized initialization, L2-SVM output layer
- Inference: Three ways: keep binary weights, use hp weights or ensemble of binary networks (sampled)
- Compute dropout+DNN Performance numbers for CIFAR + SVHN?
- Fig 1: feature coadaptation comparison
- Ref. [39] about binary DNN: Fixed-Point feedforward deep Neural...retrains network!
- 2/3 of all muls are due to forward/backward prop
- No demonstration of power saving/efficiency gains

Neural Networks with few Multiplications

- Muls again in focus (“Multiplier light networks”)
- Binarization of weights + quantization (!) of activations (but only in backprop), activations in forward phase are not quantized!
- No quantization of BN or ADAM
- Convert muls in backprop to bitshifts by pow2 quantization of activations
- Ref backprop without muls 1999
- First mention of random number generation cost
- Ternary connect -1,0,+1, sampling scheme similar to lp reference!
- how does hp error signal back prop? how about hp input?
- Most muls in weight updates ($2MN+3M$ muls at least)
- Mention muls incurred by batch norm! Cost of division!
- Table 2 analysis of total num of muls
- Compares only to ordinary SGD (but ok they use the same for the QNN)
- QNN converges slower but better in the end fig 1
- They don't explain the bit shift operations properly
- non-uniform distribution of activations
- Explanation of regularization by lp weights - VERY similar to entropy SGD ideas! large-basin solution, small description length, Additional reference (Neelakantan: Adding gradient noise)

Quantized Neural Networks

- Very lp activations + 6bit gradients (enabling bit-wise ops)
- first to actually consider gradient of det./stoch. binary units.
- first experiments on RNNs + ImageNet
- Extensive references (e.g. network compression for inference)
- MAC ops replaced by XNOR and popcount ops
- CNNs benefit mainly from activation quantization (large neuron to weight ratio)
- Binary kernel repetition (worth mentioning?)
- Abandon stochastic binarization (only in activations at train time)
- NEW: Gradient estimators (clipped straight through) due to binary activations. But: weights are binarized as well, but still they take gradients w.r.t. weights as if they were continuous...
- Shaky argument about independence of binarizations for the “derivation” of the clipped ST estimator. But intuitively ok (eq. 6)
- Address cost of BN, which is particularly large for CNNs + inverse sqrt
- How does the approximate shift work in case of negative argument?
- Make screenshot of “very unoptimized” baseline gpu kernel (theano repo), shift based BN+ADAM are in Torch repo

- Algorithm 1: input binarization? activation functions missing - no, they use det./stoch. binary units!
- Need to disentangle back-prop equations...
- What about multiplications in BackBatchNorm? No mention.
- Input handling: 8bit fixed point , check algorithm 4...
- More than one bit: Quantization+XNORpopcount, Ref to DoReFa net and logarithmic data representation papers.
- Extensions:
 - Approximate BackBatchNorm as well (if not already done)
 - Use non-uniform quantization scheme

ZipML Framework

- They apply (non-uniform) quantization scheme only to input data (?)
- How would you apply such an adaptive scheme to weights, activations, gradients?

XNOR-Net

- Their comparison to BNN on ImageNet is really strange, did they implement it badly? Because back then BNN did not yet have numbers for imagenet, but they claim it performs worse than BNN. Only later Bengio published (better!) numbers than XNOR.
- Memory/op count for AlexNet (compare DNN and CNN weight/activation counts)
- They consider binary weights+activations and only binary weights.
- Mention several network compression approaches (on pretrained networks)
- Quite extensive review of related techniques.
-
-